



AprèsSQI: Extra Fast Verification for SQIsign Using Extension-Field Signing

Maria Corte-Real Santos^{1(✉)}, Jonathan Komada Eriksen², Michael Meyer³,
and Krijn Reijnders⁴

¹ University College London, London, UK

`maria.santos.20@ucl.ac.uk`

² Norwegian University of Science and Technology, Trondheim, Norway

`jonathan.k.eriksen@ntnu.no`

³ University of Regensburg, Regensburg, Germany

`michael@random-oracles.org`

⁴ Radboud University, Nijmegen, The Netherlands

`krijn@cs.ru.nl`

Abstract. We optimise the verification of the SQIsign signature scheme. By using field extensions in the signing procedure, we are able to significantly increase the amount of available rational 2-power torsion in verification, which achieves a significant speed-up. This, moreover, allows several other speed-ups on the level of curve arithmetic. We show that the synergy between these high-level and low-level improvements gives significant improvements, making verification 2.07 times faster, or up to 3.41 times when using size-speed trade-offs, compared to the state of the art, without majorly degrading the performance of signing.

Keywords: post-quantum cryptography · isogenies · SQIsign · verification

1 Introduction

Research has shown that large-scale quantum computers will break current public-key cryptography, such as RSA or ECC, whose security relies on the hardness of integer factorization or the discrete logarithm, respectively [29]. Post-quantum cryptography seeks to thwart the threat of quantum computers by developing cryptographic primitives based on alternative mathematical problems that cannot be solved efficiently by quantum computers. In recent years, lattice-based cryptography has developed successful post-quantum schemes for essential primitives such as key encapsulation mechanisms (KEMs) and digital signatures that will be standardized by NIST. Lattice-based signatures are able

Author list in alphabetical order; see <https://www.ams.org/profession/leaders/CultureStatement04.pdf>. This work has been supported by UK EPSRC grant EP/S022503/1 and by the German Federal Ministry of Education and Research (BMBF) under the project 6G-RIC (ID 16KISK033).

The full version of this paper is available at <https://ia.cr/2023/1559>.

© International Association for Cryptologic Research 2024

M. Joye and G. Leander (Eds.): EUROCRYPT 2024, LNCS 14651, pp. 63–93, 2024.

https://doi.org/10.1007/978-3-031-58716-0_3

to provide fast signing and verification, but have to resort to larger key and signature sizes than were previously acceptable in pre-quantum signatures. For applications where the amount of data transmitted is crucial, these lattice-based schemes may not be a practical option. NIST is therefore looking for other digital signature schemes with properties such as smaller combined public key and signature sizes to ensure a smooth transition to a post-quantum world [32].

A potential solution to this problem is provided by the sole isogeny-based candidate in NIST’s new call for signatures – **SQIsign** [18] – as it is currently the candidate that comes closest to the data sizes transmitted (i.e. the combined size of the signature and the public key) in pre-quantum elliptic curve signatures [23, 24]. **SQIsign** is most interesting in scenarios that require small signature sizes and fast verification, particularly in those applications where the performance of signing is not the main concern. A few common examples include long-term signatures, specifically public-key certificates, code updates for small devices, authenticated communication with embedded devices or other microcontrollers that solely run verification, and smart cards. For such use cases it is imperative to bring down the cost of verification as much as possible.

Performance Bottlenecks in SQIsign. The bottleneck of verification in **SQIsign** is the computation of an isogeny of fixed degree 2^e with $e \approx (15/4) \log(p)$, where p denotes the prime one is working over, e.g. $\log(p) \approx 256$ for NIST Level I security. However, the rational 2-power torsion, from here on denoted as the 2^\bullet -torsion, is limited, since we work with supersingular elliptic curves over \mathbb{F}_{p^2} of order $(p+1)^2$ and $(p-1)^2$. This sets a theoretical limit of $2^{\log p}$ for the 2^\bullet -torsion. Therefore, the verifier needs to perform several *blocks* of degree 2^\bullet to complete the full 2^e -isogeny, where each of these blocks involves costly steps such as computing a 2^\bullet -torsion basis or isogeny kernel generator. Hence, in general, a smaller number of blocks improves the performance of verification.

On the other hand, the bottleneck in signing is the computation of several T -isogenies for odd smooth $T \approx p^{5/4}$. Current implementations of **SQIsign** therefore require $T \mid (p-1)(p+1)$, such that \mathbb{F}_{p^2} -rational points are available for efficient T -isogeny computations. The performance of this step is dominated by the smoothness of T , i.e., its largest prime factor.

While this additional divisibility requirement theoretically limits the maximal 2^\bullet -torsion to roughly $p^{3/4}$, current techniques for finding **SQIsign**-friendly primes suggest that achieving this with acceptable smoothness of T is infeasible [7, 8, 11, 15, 18]. In particular, the NIST submission of **SQIsign** uses a prime with rational 2^{75} -torsion and 1973 as largest factor of T . Since $e \approx (15/4) \cdot 256 = 960$, this means that the verifier has to perform $\lceil e/75 \rceil = 13$ costly isogeny blocks. Increasing the 2^\bullet -torsion further is difficult as it decreases the probability of finding a smooth and large enough T for current implementations of **SQIsign**.

Our Contributions. In this work, we deploy a range of techniques to increase the 2^\bullet -torsion and push the **SQIsign** verification cost far below the state of the art. Alongside these technical contributions, we aim to give an accessible description

of **SQIsign**, focusing primarily on verification, which solely uses elliptic curves and isogenies and does not require knowledge of quaternion algebras.

Even though we target faster verification, our main contribution is signing with field extensions. From this, we get a much weaker requirement on the prime p , which in turn enables us to increase the size of the 2^\bullet -torsion.

Focusing on NIST Level I security, we study the range of possible 2^\bullet -torsion to its theoretical maximum, and measure how its size correlates to verification time through an implementation that uses an equivalent to the number of field multiplications as cost metric. Compared to the state of the art, increasing the 2^\bullet -torsion alone makes verification almost 1.7 times faster. Further, we implement the new signing procedure as proof-of-concept in SageMath and show that signing times when signing with field extensions are in the same order of magnitude as when signing only using operations in \mathbb{F}_{p^2} .

For verification, in addition to implementing some known general techniques for improvements compared to the reference implementation provided in the NIST submission of **SQIsign**, we show that increasing the 2^\bullet -torsion also opens up a range of optimisations that were previously not possible. For instance, large 2^\bullet -torsion allows for an improved challenge-isogeny computation and improved basis and kernel generation. Furthermore, we show that size-speed trade-offs as first proposed by De Feo, Kohel, Leroux, Petit, and Wesolowski [18] become especially worthwhile for large 2^\bullet -torsion. When pushing the 2^\bullet -torsion to its theoretical maximum, this even allows for uncompressed signatures, leading to significant speed-ups at the cost of roughly doubling the signature sizes.

For two specific primes with varying values of 2^\bullet -torsion, we combine all these speed-ups, and measure the performance of verification. Compared to the implementation of the **SQIsign** NIST submission [8], we reach a speed-up up to a factor 2.70 at NIST Level I when keeping the signature size of 177 bytes. When using our size-speed trade-offs, we reach a speed-up by a factor 3.11 for signatures of 187 bytes, or a factor 4.46 for uncompressed signatures of 322 bytes. Compared to the state of the art [26], these speed-ups are factors 2.07, 2.38 and 3.41 respectively.

Related Work. De Feo, Kohel, Leroux, Petit, and Wesolowski [18] published the first **SQIsign** implementation, superseded by the work of De Feo, Leroux, Longa, and Wesolowski [19]. Subsequently, Lin, Wang, Xu, and Zhao [26] introduced several improvements for this implementation. The NIST submission of **SQIsign** [8] features a new implementation that does not rely on any external libraries. Since this is the latest and best documented implementation, we will use it as a baseline for performance comparison, and refer to it as **SQIsign** (NIST). Since the implementation from [26] is not publicly available, we included their main improvement for verification in **SQIsign** (NIST), and refer to this as **SQIsign** (LWXZ).

Dartois, Leroux, Robert, and Wesolowski [16] recently introduced **SQIsignHD**, which massively improves the signing time in **SQIsign**, in addition to a number of other benefits, but at the cost of a still unknown slowdown in verification.

This could make **SQlsignHD** an interesting candidate for applications that prioritise the combined cost of signing and verification over the sole cost of verification.

Recent work by Eriksen, Panny, Sotáková, and Veroni [21] explored the feasibility of computing the Deuring correspondence (see Sect. 2.2) for *general* primes p via using higher extension fields. We apply the same techniques and tailor them to *specialised* primes for use in the signing procedure of **SQlsign**.

Overview. The rest of the paper is organised as follows. Section 2 recalls the necessary background, including a high-level overview of **SQlsign**. Section 3 describes how using field extensions in signing affects the cost and relaxes requirements on the prime. Section 4 analyses how the size of the 2^\bullet -torsion correlates to verification time. Section 5 presents optimisations enabled by the increased 2^\bullet -torsion, while Sect. 6 gives further optimisations enabled by increased signature sizes. Finally, Sect. 7 gives some example parameters, and measures their performance compared to the state of the art.

Availability of Software. We make our Python and SageMath software publicly available under the MIT licence at

<https://github.com/TheSICQ/APresSQL>.

2 Preliminaries

Throughout this paper, p denotes a prime number and \mathbb{F}_{p^k} the finite field with p^k elements, where $k \in \mathbb{Z}_{>0}$.

2.1 Elliptic Curves and Their Endomorphism Rings

We first give the necessary geometric background to understand the **SQlsign** signature scheme. For a more general exposition we refer to Silverman [31].

Isogenies. An isogeny $\varphi : E_1 \rightarrow E_2$ between two elliptic curves E_1, E_2 is a non-constant morphism that sends the identity of E_1 to the identity of E_2 . The degree $d = \deg(\varphi)$ of an isogeny is its degree as a rational map. If the degree d of an isogeny φ has the prime factorisation $d = \prod_{i=1}^n \ell_i^{e_i}$, we can decompose φ into the composition of e_i isogenies of degree ℓ_i for $i = 1$ to n . For every isogeny $\varphi : E_1 \rightarrow E_2$, there is a (unique) *dual* isogeny $\hat{\varphi} : E_2 \rightarrow E_1$ that satisfies $\hat{\varphi} \circ \varphi = [\deg(\varphi)]$, the multiplication-by- $\deg(\varphi)$ map on E_1 . Similarly, $\varphi \circ \hat{\varphi}$ is the multiplication-by- $\deg(\varphi)$ map on E_2 .

A *separable* isogeny is described, up to isomorphism, by its kernel, a group of order d . Given a kernel G of prime order d , we can compute the corresponding isogeny $\phi : E \rightarrow E/G$ using Vélú's formulas [34] in $\tilde{O}(d)$. Leroux, and Smith [5] showed that this can be asymptotically reduced to $\tilde{O}(\sqrt{d})$ using $\sqrt{\text{elu}}$ formulas. In Sect. 2.5, we return to the topic of computing isogenies and give a more detailed discussion.

Endomorphism Rings. An isogeny from a curve E to itself is called an *endomorphism*. For example, for any integer n , the multiplication-by- n map is an endomorphism. Another, not necessarily distinct, example for elliptic curves defined over \mathbb{F}_q is the *Frobenius endomorphism* $\pi : (x, y) \mapsto (x^q, y^q)$.

The set of endomorphisms $\text{End}(E)$ of an elliptic curve E forms a ring under (pointwise) addition and composition of isogenies. The endomorphism ring of E/\mathbb{F}_p is either isomorphic to an imaginary quadratic order, or to a maximal order in a quaternion algebra ramified at p and ∞ (which will be defined in Sect. 2.2). In the latter case, we say that E is *supersingular*, and from this point forward, E will denote a supersingular curve.

Supersingular Elliptic Curves and Their Isomorphism Classes. We will mostly consider supersingular elliptic curves *up to isomorphism*, and thus work with isomorphism classes of these curves. Throughout, we will exploit the fact that every isomorphism class of supersingular curves has a model over \mathbb{F}_{p^2} , such that the p^2 -power Frobenius π is equal to the multiplication-by- $(-p)$ map. Such curves E are \mathbb{F}_{p^2} -isogenous to curves defined over \mathbb{F}_p , and satisfy

$$E(\mathbb{F}_{p^{2k}}) = E[p^k - (-1)^k] \cong \mathbb{Z}/(p^k - (-1)^k)\mathbb{Z} \oplus \mathbb{Z}/(p^k - (-1)^k)\mathbb{Z}, \quad (1)$$

while their quadratic twist over $\mathbb{F}_{p^{2k}}$, which we will denote E_k^t , satisfies

$$E_k^t(\mathbb{F}_{p^{2k}}) = E[p^k + (-1)^k] \cong \mathbb{Z}/(p^k + (-1)^k)\mathbb{Z} \oplus \mathbb{Z}/(p^k + (-1)^k)\mathbb{Z}. \quad (2)$$

For such curves, for any positive integer $N \mid p^k \pm 1$, the full N -torsion group $E[N]$ is defined over $\mathbb{F}_{p^{2k}}$, either on the curve itself, or on its twist.

The Isogeny Problem. The fundamental hard problem underlying the security of all isogeny-based primitives is the following: given two elliptic curves E_1, E_2 defined over \mathbb{F}_{p^2} find an isogeny $\phi : E_1 \rightarrow E_2$. The best classical attack against this problem is due to Delfs and Galbraith [20] which runs in time $\tilde{O}(\sqrt{p})$, and quantum attack due to Biasse, Jao, and Sankar [6] that runs in $\tilde{O}(\sqrt[4]{p})$. A related problem, which will be useful in the context of **SQIsign**, is the *endomorphism ring problem*, which asks, given a supersingular elliptic curve E/\mathbb{F}_{p^2} , to find the endomorphism ring $\text{End}(E)$. Wesolowski [36] showed that this is equivalent to the isogeny problem under reductions of polynomial expected time, assuming the generalised Riemann hypothesis, and further, Page and Wesolowski [28] recently showed that the endomorphism ring problem is equivalent to the problem of computing one endomorphism.

2.2 Quaternion Algebras and the Deuring Correspondence

We give the necessary arithmetic background to understand the signing procedure of **SQIsign** at a high level.¹ The heart of the signing procedure in **SQIsign**

¹ This section is only necessary for Sect. 2.3 and Sect. 3, as all other sections are concerned only with **SQIsign** verification, which will only use well-known isogeny terminology. In contrast, signing heavily relies on the arithmetic of quaternion algebras.

lies in the Deuring correspondence, which connects the geometric world of super-singular curves from Sect. 2.1 to the arithmetic world of quaternion algebras. For more details on quaternion algebras, we refer to Voight [35].

Quaternion Algebras, Orders and Ideals. Quaternion algebras are four-dimensional \mathbb{Q} -algebras, generated by four elements $1, i, j, k$ following certain multiplication rules. For **SQSign**, we work in the quaternion algebra *ramified* at p and ∞ . When $p \equiv 3 \pmod{4}$, one representation of such a quaternion algebra is given by $\mathcal{B}_{p,\infty} = \mathbb{Q} + i\mathbb{Q} + j\mathbb{Q} + k\mathbb{Q}$ with multiplication rules $i^2 = -1$, $j^2 = -p$, $ij = -ji = k$. For an element $\alpha = x + yi + zj + wk \in \mathcal{B}_{p,\infty}$ with $x, y, z, w \in \mathbb{Q}$, we define its *conjugate* to be $\bar{\alpha} = x - yi - zj - wk$, and its *reduced norm* to be $n(\alpha) = \alpha\bar{\alpha}$.

We are mainly interested in *lattices* in $\mathcal{B}_{p,\infty}$, defined as full-rank \mathbb{Z} -modules contained in $\mathcal{B}_{p,\infty}$, i.e., abelian groups of the form $\alpha_1\mathbb{Z} + \alpha_2\mathbb{Z} + \alpha_3\mathbb{Z} + \alpha_4\mathbb{Z}$, where $\alpha_1, \alpha_2, \alpha_3, \alpha_4 \in \mathcal{B}_{p,\infty}$ are linearly independent. If a lattice $\mathcal{O} \subset \mathcal{B}_{p,\infty}$ is also a subring of $\mathcal{B}_{p,\infty}$, i.e., it contains 1 and is closed under multiplication, then \mathcal{O} is called an *order*. Orders that are not strictly contained in any other order are called *maximal* orders. From this point on, we only consider maximal orders.

A lattice I that is closed under multiplication by an order \mathcal{O} on the left is called a *left* (resp. *right*) \mathcal{O} -*ideal*. We refer to \mathcal{O} as the left (resp. right) order of I . When \mathcal{O} is the left order of I and \mathcal{O}' the right order of I , we define I to be a *connecting* $(\mathcal{O}, \mathcal{O}')$ -*ideal*.² A left \mathcal{O} -ideal I that is also contained in \mathcal{O} is called an *integral* ideal. From this point on, we only deal with integral left ideals and simply refer to them as ideals.

The *norm* of an ideal I is the greatest common divisor of the reduced norms of the elements of I , whereas the *conjugate* \bar{I} of an ideal I is the ideal consisting of the conjugates of the elements of I . Two ideals I and J are said to be equivalent if $I = J\alpha$ for some $\alpha \in \mathcal{B}_{p,\infty}^\times$ and is denoted $I \sim J$. Equivalent ideals have equal left orders and isomorphic right orders.

The Deuring Correspondence. Given an elliptic curve E with $\text{End}(E) \cong \mathcal{O}$, there is a one-to-one correspondence between separable isogenies from E and left \mathcal{O} -ideals I of norm coprime to p .

Given an isogeny φ , we denote the corresponding ideal I_φ , and conversely, given an ideal I , we denote the corresponding isogeny φ_I . The Deuring correspondence acts like a dictionary: a given isogeny $\varphi : E \rightarrow E'$ corresponds to an ideal I_φ with left order $\mathcal{O} \cong \text{End}(E)$ and right order $\mathcal{O}' \cong \text{End}(E')$. Furthermore, the degree of φ is equal to the norm of I_φ and the dual isogeny $\hat{\varphi}$ corresponds to the conjugate $\bar{I}_\varphi = I_{\hat{\varphi}}$. Equivalent ideals I, J have isomorphic right orders and the corresponding isogenies φ_I, φ_J have isomorphic codomains.

The (Generalised) KLPT Algorithm. The KLPT algorithm, introduced by Kohel, Lauter, Petit, and Tignol [25], is a purely quaternionic algorithm,

² Note that \mathcal{O} and \mathcal{O}' need not be distinct.

but has seen a variety of applications in isogeny-based cryptography due to the Deuring correspondence. Given an ideal I , KLPT finds an equivalent ideal J of prescribed norm. The drawback is that the norm of the output J will be comparatively large.

For example, the KLPT algorithm is used to compute isogenies between two curves of known endomorphism ring. Given two maximal orders $\mathcal{O}, \mathcal{O}'$, translating the standard choice³ of connecting ideal I to its corresponding isogeny is hard. However, by processing I through KLPT first, we can find an equivalent ideal J of smooth norm, allowing us to compute φ_J . This is essential for computing the response in **SQIsign**.

The original KLPT algorithm only works for \mathcal{O}_0 -ideals, where \mathcal{O}_0 is a maximal order of a special form.⁴ This was generalised by De Feo, Kohel, Leroux, Petit, and Wesolowski [18] to work for arbitrary orders \mathcal{O} , albeit at the cost of an even larger norm bound for the output. Note that **SQIsign** utilizes both versions.

2.3 SQIsign

Next, we give a high-level description of signing and verification in **SQIsign**. **SQIsign** is a signature scheme based on an underlying Sigma protocol that proves knowledge of a *secret* (non-scalar) endomorphism $\alpha \in \text{End}(E_A)$ for some *public* curve E_A . At its core, the prover shows this knowledge by being able to compute an isogeny φ from E_A to some random curve E_2 .

A high-level description of the **SQIsign** Sigma protocol is given below.

Setup: Fix a prime number p and supersingular elliptic curve E_0/\mathbb{F}_{p^2} with known endomorphism ring.

Key generation: Compute a secret key $\varphi_A : E_0 \rightarrow E_A$, giving the prover knowledge of $\text{End}(E_A)$, with corresponding public verification key E_A .

Commit: The prover generates a random commitment isogeny $\varphi_{\text{com}} : E_0 \rightarrow E_1$, and sends E_1 to the verifier.

Challenge: The verifier computes a random challenge isogeny $\varphi_{\text{chall}} : E_1 \rightarrow E_2$, and sends φ_{chall} to the prover.

Response: The prover uses the knowledge of φ_{com} and φ_{chall} to compute $\text{End}(E_2)$, allowing the prover to compute the response isogeny $\varphi_{\text{resp}} : E_A \rightarrow E_2$, by translating an ideal computed using the generalised KLPT algorithm, as described at the end of Sect. 2.2.

The verifier needs to check that φ_{resp} is an isogeny from E_A to E_2 .⁵ Assuming the hardness of the endomorphism ring problem, the protocol is sound: if the prover is able to respond to two different challenges $\varphi_{\text{chall}}, \varphi'_{\text{chall}}$ with φ_{resp} and φ'_{resp} , the prover knows an endomorphism of the public key E_A , namely

³ $I = N\mathcal{O}\mathcal{O}'$, where N is the smallest integer making I integral.

⁴ Specifically, it only works for special, p -extremal orders. An example of such an order when $p \equiv 3 \pmod{4}$ is $\text{End}(E_0)$ where $j(E_0) = 1728$.

⁵ Additionally, $\widehat{\varphi_{\text{chall}}} \circ \varphi_{\text{resp}}$ needs to be cyclic. Observe that otherwise, the soundness proof might return a scalar endomorphism.

$\widehat{\varphi'_{\text{resp}}} \circ \varphi'_{\text{chall}} \circ \widehat{\varphi_{\text{chall}}} \circ \varphi_{\text{resp}}$. Proving zero-knowledge is harder and relies on the output distribution of the generalised KLPT algorithm. Note that KLPT is needed for computing the response:⁶ while setting $\varphi_{\text{resp}} = \varphi_{\text{chall}} \circ \varphi_{\text{com}} \circ \widehat{\varphi_A}$ gives an isogeny from E_A to E_2 , this leaks the secret φ_A .⁷ For a further discussion on zero-knowledge, we refer to the original **SQIsign** articles [18, 19].

Remark 1. The best-known attacks against **SQIsign** are the generic attacks against the endomorphism ring problem. As discussed in Sect. 2.1, their run time depends only on the size of p and, with high probability, do not recover the original secret isogeny, but rather a different isogeny between the same curves. Therefore, their complexity should be unaffected by the changes we introduce to the **SQIsign** protocol in Sect. 3, as for these attacks it does not matter whether the original secret isogeny had kernel points defined over a larger extension field. In short, the changes in this work do not affect the security of **SQIsign**.

Verification. Using the Fiat–Shamir heuristic, the **SQIsign** Sigma protocol is transformed into a signature scheme. This means that a signature on the message msg is of the form $\sigma = (\varphi_{\text{resp}}, \text{msg}, E_1)$. For efficiency, φ_{resp} is compressed, and E_1 is replaced by a description of φ_{chall} . Thus, given the signature σ and public key E_A , the verifier recomputes the response isogeny $\varphi_{\text{resp}} : E_A \rightarrow E_2$ and the (dual of the) challenge isogeny $\widehat{\varphi_{\text{chall}}} : E_2 \rightarrow E_1$, and then verifies that the hash $H(\text{msg}, E_1)$ indeed generates φ_{chall} .

The isogeny φ_{resp} is of degree 2^e with $e = \lceil \frac{15}{4} \log(p) \rceil + 25$ [8, Sect. 7.2.3], where 2^e corresponds to the output size of the generalised KLPT algorithm. The bottleneck in verification is the (re)computation of φ_{resp} in $\lceil e/f \rceil$ steps of size 2^f . Accelerating this will be the focus of this paper.

2.4 SQIsign-Friendly Primes

Next, we give more details on the parameter requirements in **SQIsign**. We refer to the original **SQIsign** works [8, 18, 19] for a detailed description of their origins.

SQIsign Prime Requirements. The main bottleneck of **SQIsign** is the computation of isogenies. Recall from Eqs. (1) and (2) that, when working with supersingular elliptic curves E/\mathbb{F}_{p^2} , we have $E(\mathbb{F}_{p^2}) = E[p \pm 1]$. Thus, to use x -only arithmetic over \mathbb{F}_{p^2} , **SQIsign** restricts to computing isogenies of *smooth* degree $N \mid (p^2 - 1)$. Finding **SQIsign**-friendly primes reduces to finding primes p , with $p^2 - 1$ divisible by a large, smooth number. More explicitly, for a security level λ , the following parameters are needed:

⁶ Alternatively, one can replace the connecting ideal with the shortest equivalent ideal, and translate it by embedding it in an isogeny between higher-dimensional abelian varieties, as shown in **SQIsignHD** [16].

⁷ Further, this is not a valid response, since the composition with $\widehat{\varphi_{\text{chall}}}$ is not cyclic.

- A prime p of bitsize $\log_2(p) \approx 2\lambda$ with $p \equiv 3 \pmod{4}$.
- The torsion group $E[2^f]$ as large as possible, that is $2^f \mid p+1$.
- A smooth odd factor $T \mid (p^2 - 1)$ of size roughly $p^{5/4}$.
- The degree of $\varphi_{\text{com}}, D_{\text{com}} \mid T$, of size roughly $2^{2\lambda} \approx p$.
- The degree of $\varphi_{\text{chall}}, D_{\text{chall}} \mid 2^f T$, of size roughly $2^\lambda \approx p^{1/2}$.
- Coprimality between D_{com} and D_{chall} .

To achieve NIST Level I, III, and V security, we set the security parameter as $\lambda = 128, 192, 256$, respectively. Concretely, this means that, for each of these security parameters, we have $\log p \approx 256, 384, 512$, and $\log T \approx 320, 480, 640$, with f as large as possible given the above restrictions. The smoothness of T directly impacts the signing time, and the problem of finding primes p with a large enough T that is reasonably smooth is difficult. We refer to recent work on this problem for techniques to find suitable primes [7, 8, 11, 15, 18, 19].

The crucial observation for this work is that T occupies space in $p^2 - 1$ that limits the size of f , hence current SQIsign primes balance the smoothness of T with the size of f .

Remark 2. SQIsign (NIST) further requires $3^g \mid p+1$ such that $D_{\text{chall}} = 2^f \cdot 3^g \geq p^{1/2}$ and $D_{\text{chall}} \mid p+1$. While this is not a strict requirement in the theoretical sense, it facilitates efficiency of computing φ_{chall} . From this point on, we ensure that this requirement is always fulfilled.

Remark 3. Since the curves E and their twists E^t satisfy

$$E(\mathbb{F}_{p^2}) \cong \mathbb{Z}/(p \pm 1)\mathbb{Z} \oplus \mathbb{Z}/(p \pm 1)\mathbb{Z},$$

and we work with both simultaneously, choosing T and f is often incorrectly described as choosing divisors of $p^2 - 1$. There is a subtle issue here: even if 2^f divides $p^2 - 1$, $E[2^f]$ may not exist as a subgroup of $\langle E(\mathbb{F}_{p^2}), \rho^{-1}(E^t(\mathbb{F}_{p^2})) \rangle \subseteq E(\mathbb{F}_{p^4})$, where $\rho : E \rightarrow E^t$ is the twisting isomorphism. While this does not usually matter in the case of SQIsign (we pick 2^f as a divisor of $p+1$, and T is odd), this becomes a problem when working over multiple extension fields. In Sect. 3.2, we make this precise and reconcile it using Theorem 1.

2.5 Computing Rational Isogenies from Irrational Generators

Finally, to facilitate signing with field extensions, we recall the techniques for computing \mathbb{F}_{p^2} -rational isogenies, i.e., isogenies defined over \mathbb{F}_{p^2} , generated by *irrational* kernel points, that is, not defined over \mathbb{F}_{p^2} . In the context of this paper, we again stress that such isogenies will only be computed by the signer; the verifier will only work with points in \mathbb{F}_{p^2} .

The main computational task of most isogeny-based cryptosystems (including SQIsign) lies in evaluating isogenies given the generators of their kernels. Explicitly, given an elliptic curve E/\mathbb{F}_q , a point $K \in E(\mathbb{F}_{q^k})$ such that $\langle K \rangle$ is defined over \mathbb{F}_q ,⁸ and a list of points (P_1, P_2, \dots, P_n) in E , we wish to compute

⁸ That is, the group $\langle K \rangle$ is closed under the action of $\text{Gal}(\bar{\mathbb{F}}_q/\mathbb{F}_q)$.

the list of points $(\varphi(P_1), \varphi(P_2), \dots, \varphi(P_n))$, where φ is the separable isogeny with $\ker \varphi = \langle K \rangle$. Since we work with curves E whose p^2 -Frobenius π is equal to the multiplication-by- $(-p)$ map (see Sect. 2.1), *every* subgroup of E is closed under the action of $\text{Gal}(\mathbb{F}_{p^2}/\mathbb{F}_{p^2})$, hence every isogeny from E can be made \mathbb{F}_{p^2} -rational, by composing with the appropriate isomorphism.

Computing Isogenies of Smooth Degree. Recall from Sect. 2.1 that the isogeny factors as a composition of small prime degree isogenies, which we compute using Vélú-style algorithms. For simplicity, for the rest of the section, we therefore assume that $\langle K \rangle$ is a subgroup of order $\ell > 2$, where ℓ is a small prime.

At the heart of these Vélú-style isogeny formulas is evaluating the kernel polynomial. Pick any subset $S \subseteq \langle K \rangle$ such that $\langle K \rangle = S \sqcup -S \sqcup \{\infty\}$. Then the kernel polynomial can be written as

$$f_S(X) = \prod_{P \in S} (X - x(P)). \quad (3)$$

Here, the generator K can be either a rational point, i.e., lying in $E(\mathbb{F}_q)$, or an irrational point, i.e., lying in $E(\mathbb{F}_{q^k})$ for $k > 1$, but whose group $\langle K \rangle$ is defined over \mathbb{F}_q . Next, we discuss how to solve the problem efficiently in the latter case.

Irrational Generators. For $K \notin E(\mathbb{F}_q)$ of order ℓ , we can speed up the computation of the kernel polynomial using the action of Frobenius. This was used in two recent works [4, 21], though the general idea was used even earlier [33].

As $\langle K \rangle$ is defined over \mathbb{F}_q , we know that the q -power Frobenius π acts as an endomorphism on $\langle K \rangle \subseteq E(\mathbb{F}_{q^k})$ and thus maps K to a multiple $[\gamma]K$ for some $\gamma \in \mathbb{Z}$. This fully determines the action on $\langle K \rangle$, i.e., $\pi|_{\langle K \rangle}$ acts as $P \mapsto [\gamma]P$ for all $P \in \langle K \rangle$. For the set S as chosen above, this action descends to an action on its x -coordinates $X_S = \{x(P) \in \mathbb{F}_{q^k} \mid P \in S\}$ and thus partitions X_S into orbits $\{x(P), x([\gamma]P), x([\gamma^2]P), \dots\}$ of size equal to the order of γ in $(\mathbb{Z}/\ell\mathbb{Z})^\times / \{1, -1\}$.

If we pick one representative $P \in S$ per orbit, and call this set of points S_0 , we can compute the kernel polynomial (3) as a product of the minimal polynomials $\mu_{x(P)}$ of the $x(P) \in \mathbb{F}_{q^k}$ for these $P \in S_0$, with each $\mu_{x(P)}$ defined over \mathbb{F}_q , as

$$f_S(X) = \prod_{P \in S_0} \mu_{x(P)}(X), \quad (4)$$

where μ_β denotes the minimal polynomial of β over \mathbb{F}_q .

To compute $f_S(\alpha)$ for $\alpha \in \mathbb{F}_q$, we only require the smaller polynomial $f_{S_0}(X)$ and compute $\text{Norm}_{\mathbb{F}_{q^k}/\mathbb{F}_q}(f_{S_0}(\alpha))$, as

$$\text{Norm}_{\mathbb{F}_{q^k}/\mathbb{F}_q}(f_{S_0}(\alpha)) = \prod_{\pi \in G} \pi(f_{S_0}(\alpha)) = \prod_{P \in S_0} \prod_{\pi \in G} (\alpha - \pi(x(P))) = \prod_{P \in S_0} \mu_{x(P)}(\alpha),$$

where $G = \text{Gal}(\mathbb{F}_{q^k}/\mathbb{F}_q)$, as per Banegas, Gilchrist, Le Dévéhat, and Smith [4]. This allows us to compute the image under f_S of x -values of points in $E(\mathbb{F}_q)$, but only works for values in \mathbb{F}_q . To evaluate $f_S(\alpha)$ for general $\alpha \in \overline{\mathbb{F}_p}$, i.e. to

compute the image of a point in $E(\overline{\mathbb{F}}_p)$, we instead use the larger polynomial $f_S(X)$, which we compute, as in Eq. (4), as a product of the minimal polynomials $\mu_{x(P)}$, where we use Shoup’s algorithm [30] to compute each $\mu_{x(P)}$ given $x(P)$. Computing $f_S(X)$ requires a total of $O(\ell k) + \tilde{O}(\ell)$ operations, with k such that each $x(P) \in \mathbb{F}_{q^k}$. Evaluation f_S at α takes $\tilde{O}(\ell k')$ operations, with k' the smallest value such that $\alpha \in \mathbb{F}_{q^{k'}}$ [21, Sect. 4.3].

Remark 4. The biggest drawback to using this technique is that $\sqrt{\ell \ln}$ is no longer effective, as we would need to work in the smallest field where both the isogeny generator and the x -value of the point we are evaluating are defined in.

3 Signing with Extension Fields

By allowing torsion T from extension fields, we enable more flexibility in choosing **SQIsign** primes p , thus enabling a larger 2^\bullet -torsion. Such torsion T requires us to compute rational isogenies with kernel points in extension fields $\mathbb{F}_{p^{2k}}$. This section describes how to adapt **SQIsign**’s signing procedure to enable such isogenies, and the increased cost this incurs. In particular, we describe two approaches for T : allowing torsion T from a particular extension field $\mathbb{F}_{p^{2k}}$, or from all extension fields $\mathbb{F}_{p^{2n}}$ for $1 \leq n \leq k$. The first approach means that we can look for T dividing an integer of bit size $\Theta(k \log p)$, and the second approach allows for $\Theta(k^2 \log p)$. In Sect. 4, we explore how increased 2^\bullet -torsion affects verification.

Throughout this section, we will reuse the notation from Sect. 2.4 to describe the various parameters related to **SQIsign**.

3.1 Changes in the Signing Procedure

Recall from Sect. 2.3 that the signing operation in **SQIsign** requires us to work with both elliptic curves and quaternion algebras, and to translate back and forth between these worlds. Note that the subroutines that work solely with objects in the quaternion algebra $\mathcal{B}_{p,\infty}$, including all operations in KLPT and its derivatives, are indifferent to what extension fields the relevant torsion groups lie in. Hence, a large part of signing is unaffected by torsion from extension fields.

In fact, the only subroutines that are affected by moving to extension fields are those relying on Algorithm 1, **IdealTolsogeny** $_D$, which translates \mathcal{O}_0 -ideals I of norm dividing D to their corresponding isogenies φ_I . **IdealTolsogeny** $_D$ is not used during verification, and is used only in the following parts of signing:

Commitment: The signer translates a random ideal of norm D_{com} to its corresponding isogeny, using one execution of **IdealTolsogeny** $_{D_{\text{com}}}$.

Response: The signer translates an ideal of norm 2^e to its corresponding isogeny, requiring $2 \cdot \lceil e/f \rceil$ executions of **IdealTolsogeny** $_T$ (see [19]).

Remark 5. We will choose parameters such that $2^f \mid p+1$ and $D_{\text{chall}} \mid p+1$, so that $E[2^f]$ and $E[D_{\text{chall}}]$ are defined over \mathbb{F}_{p^2} . As a result, the verifier only works in \mathbb{F}_{p^2} and the added complexity of extension fields applies only to the signer.

Algorithm 1. $\text{IdealTolsogeny}_D(I)$ **Input:** I a left \mathcal{O}_0 -ideal of norm dividing D **Output:** φ_I

- 1: Compute α such that $I = \mathcal{O}_0\langle\alpha, \text{nrd}(I)\rangle$
- 2: Let $\mathbf{A} = [1, i, \frac{i+j}{2}, \frac{1+k}{2}]$ denote a basis of \mathcal{O}_0
- 3: Compute $\mathbf{v}_{\bar{\alpha}} := [x_1, x_2, x_3, x_4]^T \in \mathbb{Z}^4$ such that $\mathbf{A}\mathbf{v}_{\bar{\alpha}} = \bar{\alpha}$
- 4: **for** $\ell^e \parallel D$ **do**
- 5: $\bar{\alpha}|_{\langle P_{\ell^e}, Q_{\ell^e} \rangle} := x_1\mathbf{I} + x_2(i|_{\langle P_{\ell^e}, Q_{\ell^e} \rangle}) + x_3(\frac{i+j}{2}|_{\langle P_{\ell^e}, Q_{\ell^e} \rangle}) + x_4(\frac{1+k}{2}|_{\langle P_{\ell^e}, Q_{\ell^e} \rangle})$
- 6: Let a, b, c, d be integers such that $\bar{\alpha}|_{\langle P_{\ell^e}, Q_{\ell^e} \rangle} = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$
- 7: $K_{\ell^e} := [a]P_{\ell^e} + [c]Q_{\ell^e}$
- 8: **if** $\text{ord}(K_{\ell^e}) < \ell^e$ **then**
- 9: $K_{\ell^e} = [b]P_{\ell^e} + [d]Q_{\ell^e}$
- 10: Set φ_I to be the isogeny generated by the points K_{ℓ^e} .
- 11: **return** φ_I

Adapting Ideal-to-Isogeny Translations to Field Extensions. To facilitate signing with field extensions, we slightly adapt IdealTolsogeny_D so that it works with prime powers separately. Note that the additional cost of this is negligible compared to the cost of computing the isogeny from the generators because finding the action of the relevant endomorphisms consists of simple linear algebra. See Algorithm 1 for details.

In Line 5 of Algorithm 1, the notation $\beta|_{\langle P_{\ell^e}, Q_{\ell^e} \rangle}$ refers to the action of an endomorphism β on a fixed basis P_{ℓ^e}, Q_{ℓ^e} of $E[\ell^e]$. This action is described by a matrix in $M_2(\mathbb{Z}/\ell^e\mathbb{Z})$. These matrices can be precomputed, hence the only operations in which the field of definition of $E[\ell^e]$ matters are the point additions in Lines 7 and 9, and isogenies generated by each K_{ℓ^e} in Line 10.

3.2 Increased Torsion Availability from Extension Fields

Next, we detail the two approaches to allow torsion groups from extension fields, which permits more flexibility in choosing the final prime p .

Working with a Single Field Extension of \mathbb{F}_{p^2} . Although the choice of solely working in \mathbb{F}_{p^2} occurs naturally,⁹ there is no reason *a priori* that this choice is optimal. Instead, we can choose to work in the field $\mathbb{F}_{p^{2k}}$. We emphasise that this does *not* affect signature sizes; the only drawback is that we now perform more expensive $\mathbb{F}_{p^{2k}}$ -operations during signing in IdealTolsogeny . The upside, however, is a relaxed prime requirement: we are no longer bound to $E[T] \subseteq \langle E(\mathbb{F}_{p^2}), \rho^{-1}(E^t(\mathbb{F}_{p^2})) \rangle$ and can instead use $E[T] \subseteq \langle E(\mathbb{F}_{p^{2k}}), \rho^{-1}(E^t(\mathbb{F}_{p^{2k}})) \rangle$.

By Eqs. (1) and (2), we have $E(\mathbb{F}_{p^{2k}}) \cong E[p^k \pm 1]$ and $E^t(\mathbb{F}_{p^{2k}}) \cong E[p^k \mp 1]$, thus we simply get

$$E[T] \subseteq E\left[\frac{p^{2k} - 1}{2}\right],$$

⁹ It is the smallest field over which every isomorphism class of supersingular elliptic curves has a model.

since $\langle E[A], E[B] \rangle = E[\text{lcm}(A, B)]$. Hence, by using torsion from $E(\mathbb{F}_{p^{2k}})$, we increase $T \mid (p^2 - 1)/2$ to $T \mid (p^{2k} - 1)/2$. This implies there are $2(k - 1) \log p$ more bits available to find T with adequate smoothness.

Working with Multiple Field Extensions of \mathbb{F}_{p^2} . Instead of fixing a single higher extension field $\mathbb{F}_{p^{2k}}$, we can also choose to work with multiple field extensions, in particular all fields $\mathbb{F}_{p^{2n}}$, where $1 \leq n \leq k$. The torsion group we can access by this relaxed requirement is described by the following definition.

Definition 1. *Let E be a supersingular elliptic curve over \mathbb{F}_{p^2} and let E_n^t denote an arbitrary quadratic twist of E over $\mathbb{F}_{p^{2n}}$ with respect to the twisting isomorphism $\rho_n : E \rightarrow E_n^t$. We define the k -available torsion of E to be the group generated by $E(\mathbb{F}_{p^{2n}})$ and $\rho_n^{-1}(E_n^t(\mathbb{F}_{p^{2n}}))$ for $1 \leq n \leq k$.*

Any point P in the k -available torsion can thus be written as a sum $P = \sum_{i=1}^k (P_i + \rho_n^{-1}(P_i^t))$ of points $P_i \in E(\mathbb{F}_{p^{2n}})$ and $P_i^t \in E_n^t(\mathbb{F}_{p^{2n}})$. Since the twisting isomorphism keeps the x -coordinate fixed, the computation of this isomorphism can be ignored when using x -only arithmetic, and we simply obtain a sum of points whose x -coordinates lie in $\mathbb{F}_{p^{2n}}$ for $1 \leq n \leq k$. This justifies the name k -available torsion, as we do not have to go beyond $\mathbb{F}_{p^{2k}}$ to do arithmetic with P by working with the summands separately.

The structure of the k -available torsion is completely captured by the following result, which we prove in the full version of our paper [10, Sect.3.2].

Theorem 1. *Let $p > 2$ be a prime, and let E/\mathbb{F}_{p^2} be a supersingular curve with $\text{tr}(\pi) = \pm 2p$, where π is the Frobenius endomorphism. Then the k -available torsion is precisely the group $E[N]$ with*

$$N = \prod_{n=1}^k \Phi_n(p^2)/2,$$

where Φ_n denotes the n -th cyclotomic polynomial.

We find that using all extension fields $\mathbb{F}_{p^{2n}}$, for $1 \leq n \leq k$, increases $T \mid p^2 - 1$ to $T \mid N$, with N as given by Theorem 1. Given that

$$\log(N) = \sum_{n=1}^k \log(\Phi_n(p^2)/2) \approx 2 \sum_{n=1}^k \phi(n) \log(p),$$

and the fact that $\sum_{n=1}^k \phi(n)$ is in the order of $\Theta(k^2)$, where ϕ denotes Euler's totient function, we find that $T \mid N$ gives roughly $k^2 \log(p)$ more bits to find T

with adequate smoothness, compared to the $\log(p)$ bits in the classical case of working over \mathbb{F}_{p^2} , and $k \log(p)$ bits in the case of working over $\mathbb{F}_{p^{2k}}$. Due to this, we only consider working in multiple field extensions from this point on.

3.3 Cost of Signing Using Extension Fields

In **SQsign**, operations over \mathbb{F}_{p^2} make up the majority of the cost during signing [19, Sect. 5.1]. Hence, we can roughly estimate the cost of signing by ignoring purely quaternionic operations, in which case the bottleneck of the signing procedure becomes running **IdealTolsogeny_T** as many times as required by the **IdealTolsogenyEichler** algorithm [19, Algorithm 5] in the response phase. In other words, we estimate the total signing cost from the following parameters:

- f , such that $2^f \mid p + 1$.
- T , the chosen torsion to work with.
- For each $\ell_i^{e_i} \mid T$, the smallest k_i such that $E[\ell_i^{e_i}]$ is defined over $\mathbb{F}_{p^{2k_i}}$.

Since Algorithm 1 works with prime powers separately, we can estimate the cost of a single execution by considering the cost per prime power.

Cost per Prime Power. For each $\ell_i^{e_i} \mid T$, let k_i denote the smallest integer so that $E[\ell_i^{e_i}] \subseteq E(\mathbb{F}_{p^{2k_i}})$, and let $M(k_i)$ denote the cost of operations in $\mathbb{F}_{p^{2k_i}}$ in terms of \mathbb{F}_{p^2} -operations. Computing the generator $K_{\ell_i^{e_i}}$ consists of a few point additions in $E[\ell_i^{e_i}]$, hence is $O(M(k) \cdot e \log \ell)$, while the cost of computing the isogeny generated by $K_{\ell_i^{e_i}}$ comes from computing e isogenies of degree ℓ at a cost of $O(\ell k) + \tilde{O}(\ell)$, using the techniques from Sect. 2.5.

To compute the whole isogeny, we need to push the remaining generators $K_{\ell_j^{e_j}}$, through this isogeny. To minimize the total cost, we pick the greedy strategy of always computing the smaller ℓ first. This bounds the cost of evaluating K_{ℓ^e} in *other* isogenies by $O(M(k) \cdot \ell)$.

Total Cost of Signing. Based on the analysis above, we let

$$\text{COST}_p(\ell_i^{e_i}) = c_1 M(k_i) e \log \ell + c_2 e \ell_i k_i + c_3 e \ell_i \log(\ell_i) + c_4 M(k_i) \ell$$

where k_i , and $M(k)$ are as before, and c_i are constants corresponding to the differences in the asymptotic complexities. Since we can estimate the total cost of executing **IdealTolsogeny_T** by summing the cost of each maximal prime power divisor of T , and observing that signing consists of executing **IdealTolsogeny_{D_{com}}** one time, and **IdealTolsogeny_T** a total of $2 \cdot \lceil e/f \rceil$ times, we get a rough cost estimate of signing as

$$\text{SIGNINGCOST}_p(T) = (2 \cdot \lceil e/f \rceil + 1) \cdot \sum_{\ell_i^{e_i} \mid T} \text{COST}_p(\ell_i^{e_i}).$$

In Sect. 7, we use this function to pick p and T minimising this cost. While this cost metric is very rough, we show that our implementation roughly matches the times predicted by this function. Further, we show that this cost metric suggests that going to extension fields gives signing times within the same order of magnitude as staying over \mathbb{F}_{p^2} , even when considering the additional benefit of using $\sqrt{\text{elu}}$ to compute isogenies in the latter case.

4 Effect of Increased 2^\bullet -Torsion on Verification

In Sect. 3, we showed that signing with extension fields gives us more flexibility in choosing the prime p , and, in particular, allows us to find primes with rational 2^f -torsion for larger f . In this section, we analyse how such an increase in 2^\bullet -torsion affects the cost of SQIsign verification, e.g., computing φ_{resp} and $\widehat{\varphi_{\text{chall}}}$, in terms of \mathbb{F}_p -multiplications,¹⁰ taking the SQIsign (NIST) implementation (with no further optimisations) as the baseline for comparison.

4.1 Detailed Description of Verification

Before giving an in-depth analysis of verification performance, we give a detailed description of how verification is executed. Recall that a SQIsign signature σ for a message msg created by a signer with secret signing key $\varphi_A : E_0 \rightarrow E_A$ proves knowledge of an endomorphism on E_A by describing an isogeny $\varphi_{\text{resp}} : E_A \rightarrow E_2$ of degree 2^e . A given message msg is hashed on E_1 to a point K_{chall} of order D_{chall} , hence represents an isogeny $\varphi_{\text{chall}} : E_1 \rightarrow E_2$. A signature is valid if the composition of φ_{resp} with $\widehat{\varphi_{\text{chall}}}$ is cyclic of degree $2^e \cdot D_{\text{chall}}$.

Thus, to verify a signature σ , the verifier must **(a)** recompute φ_{resp} , **(b)** compute the dual of φ_{chall} , to confirm that both are well-formed, and finally **(c)** recompute the hash of the message msg to confirm the validity of the signature.

In SQIsign, the size of the sample space for φ_{chall} impacts soundness, a key security property for signature schemes. In SQIsign (NIST), to obtain negligible soundness error (in the security parameter λ) the message is hashed to an isogeny of degree $D_{\text{chall}} = 2^f \cdot 3^g$ so that the size of cyclic isogenies of degree D_{chall} is larger than 2^λ . In contrast, when $f \geq \lambda$, we can simply set $D_{\text{chall}} = 2^\lambda$.

The signature σ consists of a compressed description of the isogenies φ_{resp} and $\widehat{\varphi_{\text{chall}}}$. For $f < \lambda$ and $D_{\text{chall}} = 2^f \cdot 3^g$ it is of the form

$$\sigma = (b, s^{(1)}, \dots, s^{(n)}, r, b_2, s_2, b_3, s_3)$$

with $s^{(j)}, s_2 \in \mathbb{Z}/2^f\mathbb{Z}$, $s_3 \in \mathbb{Z}/3^g\mathbb{Z}$, $r \in \mathbb{Z}/2^f 3^g\mathbb{Z}$, and $b, b_2, b_3 \in \{0, 1\}$. If $f \geq \lambda$, we set $D_{\text{chall}} = 2^f$ and have $s_2 \in \mathbb{Z}/2^\lambda\mathbb{Z}$ and $r \in \mathbb{Z}/2^f\mathbb{Z}$, while b_3, s_3 are omitted. Algorithmically, the verification process mostly requires three subroutines.

FindBasis. Given a curve E , find a deterministic basis (P, Q) of $E[2^f]$.

¹⁰ As standard, we denote multiplications by **M**, squarings by **S**, and additions by **a**.

FindKernel. Given a curve E with basis (P, Q) for $E[2^f]$ and $s \in \mathbb{Z}/2^f\mathbb{Z}$, compute the kernel generator $K = P + [s]Q$.

Computelsogeny. Given a curve E and a kernel generator K , compute the isogeny $\varphi : E \rightarrow E/\langle K \rangle$ and $\varphi(Q)$ for some $Q \in E$.

Below we detail each of the three verification steps **(a)**–**(c)**.

Step (a). Computing φ_{resp} is split up into $n - 1$ blocks $\varphi^{(j)} : E^{(j)} \rightarrow E^{(j+1)}$ of size 2^f , and a last block of size 2^{f_0} , where $f_0 = e - (n - 1) \cdot f$. For every $\varphi^{(j)}$, the kernel $\langle K^{(j)} \rangle$ is given by the generator $K^{(j)} = P^{(j)} + [s^{(j)}]Q^{(j)}$ for a deterministic basis $(P^{(j)}, Q^{(j)})$ of $E^{(j)}[2^f]$.

In the first block, after sampling $(P^{(1)}, Q^{(1)})$ via **FindBasis**, the bit b indicates whether $P^{(1)}$ and $Q^{(1)}$ have to be swapped before running **FindKernel**. For the following blocks, the verifier pushes $Q^{(j)}$ through the isogeny $\varphi^{(j)}$ to get a point $Q^{(j+1)} \leftarrow \varphi^{(j)}(Q^{(j)})$ on $E^{(j+1)}$ of order 2^f above $(0, 0)$.¹¹ Hence, for $j > 1$ **FindBasis** only needs to find a suitable point $P^{(j)}$ to complete the basis $(P^{(j)}, Q^{(j)})$. Furthermore, $K^{(j)}$ is never above $(0, 0)$ for $j > 1$, which ensures cyclicity when composing $\varphi^{(j)}$ with $\varphi^{(j-1)}$. In all cases we use $s^{(j)}$ from σ to compute the kernel generator $K^{(j)}$ via **FindKernel** and $\varphi^{(j)}$ via **Computelsogeny**.

The last block of degree 2^{f_0} uses $Q^{(n)} \leftarrow [2^{f-f_0}]\varphi^{(n-1)}(Q^{(n-1)})$ and samples another point $P^{(n)}$ as basis of $E^{(n)}[2^{f_0}]$. In the following, we will often assume $f_0 = f$ for the sake of simplicity.¹²

Step (b). Computing $\widehat{\varphi_{\text{chall}}}$ requires a single isogeny of smooth degree $D_{\text{chall}} \approx 2^\lambda$. For the primes given in **SQIsign** (NIST), we have $E_2[D_{\text{chall}}] \subseteq E_2(\mathbb{F}_{p^2})$. Thus, we compute φ_{chall} by deterministically computing a basis (P, Q) for $E_2[D_{\text{chall}}]$ and finding the kernel $\langle K \rangle$ for $\widehat{\varphi_{\text{chall}}} : E_2 \rightarrow E_1$. For $f < \lambda$, we have $D_{\text{chall}} = 2^f \cdot 3^g$, and split this process into two parts.

Given the basis (P, Q) for $E_2[D_{\text{chall}}]$, we compute $(P_2, Q_2) = ([3^g]P, [3^g]Q)$ as basis of $E_2[2^f]$, and use $K_2 = P_2 + [s_2]Q_2$, where b_2 indicates whether P_2 and Q_2 have to be swapped prior to computing K_2 . We compute $\varphi_2 : E_2 \rightarrow E'_2$ with kernel $\langle K_2 \rangle$, and $P_3 = [2^f]\varphi_2(P)$ and $Q_3 = [2^f]\varphi_2(Q)$ form a basis of $E'_2[3^g]$. Then b_3 indicates a potential swap of P_3 and Q_3 , while $K_3 = P_3 + [s_3]Q_3$ is the kernel generator of the isogeny $\varphi_3 : E'_2 \rightarrow E_1$. Thus, we have $\widehat{\varphi_{\text{chall}}} = \varphi_3 \circ \varphi_2$. If $f \geq \lambda$, we require only the first step.

We furthermore verify that the composition of φ_{resp} and $\widehat{\varphi_{\text{chall}}}$ is cyclic, by checking that the first 2-isogeny step of φ_2 does not revert the last 2-isogeny step of $\varphi^{(n)}$. This guarantees that $\widehat{\varphi_{\text{chall}}} \circ \varphi_{\text{resp}}$ is non-backtracking, hence cyclic.

Step (c). On E_1 , the verifier uses the point $Q' \leftarrow \widehat{\varphi_{\text{chall}}}(Q')$, where Q' is some (deterministically generated) point, linearly independent from the generator of $\widehat{\varphi_{\text{chall}}}$, and r (given in σ) to compute $[r]Q'$, and checks if $[r]Q'$ matches the hashed point $K_{\text{chall}} = H(\text{msg}, E_1)$ with hash function H .

¹¹ A point P is said to be *above* a point R if $[k]P = R$ for some $k \in \mathbb{N}$.

¹² In contrast to earlier versions, **SQIsign** (NIST) fixes $f_0 = f$. However, our analysis benefits from allowing $f_0 < f$.

4.2 Impact of Large f on Verification

The techniques of Sect. 3 extend the possible range of f to any size below $\log(p)$. This gives two benefits to the cost of verification, especially when $f \geq \lambda$.

Number of Blocks in φ_{resp} . The larger f is, the fewer blocks of size 2^f are performed in **Step (a)**. Per block, the dominating part of the cost are **FindBasis** and **FindKernel** as we first need to complete the torsion basis $(P^{(j)}, Q^{(j)})$ for $E^{(j)}[2^f]$ (given $Q^{(j)}$ if $j > 1$), followed by computing $K^{(j)} = P^{(j)} + [s^{(j)}]Q^{(j)}$. By minimizing the number of blocks n , we minimize the amount of times we perform **FindBasis** and **FindKernel**, and the cost of each individual **FindKernel** only mildly increases, as $s^{(j)}$ increases in size. The overall cost of **Computelsogeny**, that is, performing the n isogenies of degree 2^f given their kernels $K^{(j)}$, only moderately increases with growing f .

We further note that larger f requires fewer T -isogeny computations for the signer, hence signing performance also benefits from smaller n .

Challenge Isogeny. When $f \geq \lambda$, we can simply set $D_{\text{chall}} = 2^\lambda$, which has two main benefits. Firstly, the cost of **FindBasis** for this step is reduced as finding a basis for $E[2^\lambda]$ is much easier than a basis search for $E[2^f \cdot 3^g]$. Secondly, the cost for **Computelsogeny** for φ_{chall} decreases as we only have to compute a chain of 2-isogenies instead of additional 3-isogenies.

4.3 Implementation and Benchmark of Cost in \mathbb{F}_p -Multiplications

To measure the influence of the size of f on the performance, we implemented **SQIsign** verification for the NIST Level I security parameter set in Python, closely following **SQIsign** (NIST). As is standard in isogeny-based schemes, we use x -only arithmetic and represent points and curve coefficients projectively. The benchmark counts \mathbb{F}_p -operations and uses a cost metric that allows us to estimate the runtime of real-world implementations for 256-bit primes $p^{(f)}$, where $p^{(f)}$ denotes a prime such that 2^f divides $p^{(f)} + 1$. We benchmark primes $p^{(f)}$ for all values $50 \leq f \leq 250$. These results serve as a baseline to which we compare the optimisations that we introduce in Sects. 5 and 6.

We briefly outline how **SQIsign** (NIST) implements the three main subroutines **FindBasis**, **FindKernel**, and **Computelsogeny**.

FindBasis. We search for points of order 2^f by sampling x -coordinates in a specified order,¹³ and check if the corresponding point P lies on E (and not on its twist E^t). We then compute $P \leftarrow [\frac{p+1}{2^f}]P$ and verify that $[2^{f-1}]P \neq \infty$. Given two points $P, Q \in E$ of order 2^f , we verify linear independence by checking that $[2^{f-1}]P \neq [2^{f-1}]Q$, and discard and re-sample the second point otherwise.

FindKernel. Given a basis (P, Q) , **FindKernel** computes $K = P + [s]Q$ via the **3ptLadder** algorithm as used in SIKE [22]. In addition to the x -coordinates x_P

¹³ **SQIsign** (NIST) fixes the sequence $x_k = 1 + k \cdot i$ with $i \in \mathbb{F}_{p_2}$ such that $i^2 = -1$ and picks the smallest k for which we find a suitable point.

and x_Q of P and Q , it requires the x -coordinate x_{P-Q} of $P - Q$. Hence, after running `FindBasis`, we further compute x_{P-Q} as described in `SQIsign` (NIST) [8].

Computelsogeny. Given a kernel generator K of order 2^f , `Computelsogeny` follows the approach of `SIKE` [22], and computes the 2^f -isogeny $\varphi^{(j)}$ as a chain of 4-isogenies for efficiency reasons. If f is odd, we further compute a single 2-isogeny. Following `SQIsign` (NIST), `Computelsogeny` proceeds as follows:

1. Compute $R = [2^{f-2}]K$ and the corresponding 4-isogeny φ with kernel $\langle R \rangle$. Note that the point $(0,0)$ might be contained in $\langle R \rangle$ for the first block in φ_{resp} , which requires a special 4-isogeny formula. Thus, we check if this is the case and call the suitable 4-isogeny function. We set $K \leftarrow \varphi(K)$.
2. If f is odd, we compute $R = [2^{f-3}]K$, the 2-isogeny φ with kernel $\langle R \rangle$, and $K \leftarrow \varphi(K)$.
3. Compute the remaining isogeny of degree $2^{f'}$ with even f' as a chain of 4-isogenies, where $(0,0)$ is guaranteed not to lie in any of the kernels.

In the last step, `SQIsign` (NIST) uses *optimal strategies* as in `SIKE` [22] to compute a chain of 4-isogenies. Naive multiplicative strategies would compute $R = [2^{f'-2j}]K$, the 4-isogeny φ with kernel $\langle R \rangle$, and $K \leftarrow \varphi(K)$ for $j = 1, \dots, f'/2$. However, this strategy is dominated by costly doublings. Instead, we can save intermediate multiples of K during the computation of $R = [2^{f'-2j}]K$, and push them through isogenies to save multiplicative effort in following iterations. Optimal strategies that determine which multiples are pushed through isogenies and minimise the cost can be found efficiently [17, 22].

We note that for $f < \lambda$ the computation of $\widehat{\varphi_{\text{chall}}}$ requires small adaptations to these algorithms to allow for finding a basis of $E[D_{\text{chall}}]$ and computing 3-isogenies. Most notably, `SQIsign` (NIST) does *not* use optimised formulas or optimal strategies for 3-isogenies from `SIKE` [22], but uses a multiplicative strategy and general odd-degree isogeny formulas [13, 27]. We slightly deviate from `SQIsign` (NIST) by implementing optimised 3-isogeny formulas, but note that the performance difference is minor and in favor of `SQIsign` (NIST).

Cost Metric. In implementations, \mathbb{F}_{p^2} -operations usually call underlying \mathbb{F}_p -operations. We follow this approach and use the total number of \mathbb{F}_p -operations in our benchmarks. As cost metric, we express these operations in terms of \mathbb{F}_p -multiplications, with $\mathbf{S} = 0.8 \cdot \mathbf{M}$, ignoring \mathbb{F}_p -additions and subtractions due to their small impact on performance. \mathbb{F}_p -inversions, \mathbb{F}_p -square roots, and Legendre symbols over \mathbb{F}_p require exponentiations by an exponent in the range of p , hence we count their cost as $\log p$ \mathbb{F}_p -multiplications. In contrast to measuring clock cycles of an optimised implementation, our cost metric eliminates the dependence on the level of optimisation of finite field arithmetic and the specific device running `SQIsign`, hence, can be considered more general.

Benchmark Results. Figure 1 shows the verification cost for the NIST Level I-sized primes $p^{(f)}$ for $50 \leq f \leq 250$, fixing $e = 975$, using our cost metric. For more efficient benchmarking, we sample random public key curves and signatures σ of the correct form instead of signatures generated by the `SQIsign` signing procedure.




Fig. 1. Cost in \mathbb{F}_p -multiplications for verification at NIST Level I security, for varying f and $p^{(f)}$, averaged over 1024 runs per prime. The green vertical lines mark $f = 75$ as used in **SQIsign** (NIST) for signing without extension fields, and $f = \lambda = 128$, beyond which we can set $D_{\text{chall}} = 2^\lambda$. The dotted graph beyond $f = 75$ is only accessible when signing with extension fields. (Color figure online)

The graph shows the improvement for $f \geq 128$. Furthermore, we can detect when the number of blocks n decreases solely from the graph (e.g. $f = 122, 140, 163, 195, 244$). The cost of sampling a 2^f -torsion basis is highly variable between different runs for the same prime, which is visible from the oscillations of the graph. The performance for odd f is worse in general due to the inefficient integration of the 2-isogeny, which explains the zigzag-shaped graph.

From the above observations, we conclude that $f \geq \lambda$ is significantly faster for verification, with local optima found at $f = 195$ and $f = 244$, due to those being (almost) exact divisors of the signing length $e = 975$.

Remark 6. The average cost of **FindBasis** differs significantly between primes p even if they share the same 2^f -torsion. This happens because **SQIsign** (NIST) finds basis points from a pre-determined sequence $[x_1, x_2, x_3, \dots]$ with $x_j \in \mathbb{F}_{p^2}$. As we will see in Sect. 5, these x_j values can not be considered random: some values x_j are certain to be above a point of order 2^f , while others are certain not to be, for any supersingular curve over p .

5 Optimisations for Verification

In this section, we show how the improvements from Sect. 3 that increase f beyond λ together with the analysis in Sect. 4 allow several other optimisations

that improve the verification time of **SQlsign** in practice. Whereas the techniques in Sect. 3 allow us to decrease the *number* of blocks, in this section, we focus on the operations occurring *within blocks*. We optimise the cost of **FindBasis**, **FindKernel** and **Computelogsogeny**.

We first analyse the properties of points that have full 2^f -torsion, and use these properties to improve **FindBasis** and **FindKernel** for general f . We then describe several techniques specifically for $f \geq \lambda$. Altogether, these optimisations significantly change the implementation of verification in comparison to **SQlsign** (NIST). We remark that the implementation of the signing procedure must be altered accordingly, as exhibited by our implementation.

Notation. As we mostly focus on the subroutines *within* a specific block $E^{(j)} \rightarrow E^{(j+1)}$, we will omit the superscripts in $E^{(j)}$, $K^{(j)}$, $P^{(j)}$, \dots and write E, K, P, \dots to simplify notation.

5.1 Basis Generation for Full 2-Power Torsion

We first give a general result on points having full 2^f -torsion that we will use throughout this section. This theorem generalises previous results [14, 26] and will set the scene for easier and more efficient basis generation for $E[2^f]$. Its proof can be found in the full version of our paper [10, Sect. 5.1].

Theorem 2. *Let $E : y^2 = (x - \lambda_1)(x - \lambda_2)(x - \lambda_3)$ be an elliptic curve over \mathbb{F}_{p^2} with $E[2^f] \subseteq E(\mathbb{F}_{p^2})$ the full 2-power torsion. Let $L_i = (\lambda_i, 0)$ denote the points of order 2 and $[2]E$ denote the image of E under $[2] : P \mapsto P + P$ so that $E \setminus [2]E$ are the points with full 2^f -torsion. Then*

$$Q \in [2]E \text{ if and only if } x_Q - \lambda_i \text{ is square for } i = 1, 2, 3.$$

More specifically, for $Q \in E \setminus [2]E$, Q is above L_i if and only if $x_Q - \lambda_i$ is square and $x_Q - \lambda_j$ is non-square for $j \neq i$.

Note that for Montgomery curves $y^2 = x^3 + Ax^2 + x = x(x - \alpha)(x - 1/\alpha)$, the theorem above tells us that non-squareness of x_Q for $Q \in E(\mathbb{F}_{p^2})$ is enough to imply Q has full 2^f -torsion and is not above $(0, 0)$ [26, Theorem 3].

Finding Points with 2^f -Torsion Above $(0, 0)$. We describe two methods to efficiently sample Q above $(0, 0)$, based on Theorem 2.

1. **Direct x sampling.** By deterministically sampling $x_Q \in \mathbb{F}_p$, we ensure that x_Q is square in \mathbb{F}_{p^2} . Hence, if Q lies on E and $x_Q - \alpha \in \mathbb{F}_{p^2}$ is non-square, where α is a root of $x^2 + Ax + 1$, then Theorem 2 ensures that $Q \in E \setminus [2]E$ and above $(0, 0)$.
2. **Smart x sampling.** We can improve this using the fact that α is always square [2, 12]. Hence, if we find $z \in \mathbb{F}_{p^2}$ such that z is square and $z - 1$ is non-square, we can choose $x_Q = z\alpha$ square and in turn $x_Q - \alpha = (z - 1)\alpha$ non-square. Again, by Theorem 2 if Q is on E , this ensures Q is above $(0, 0)$ and contains full 2^f -torsion. Hence, we prepare a list $[z_1, z_2, \dots]$ of such values z for a given prime, and try $x_j = z_j\alpha$ until x_j is on E .

Both methods require computing α , dominated by one \mathbb{F}_{p^2} -square root. Direct sampling computes a Legendre symbol of $x^3 + Ax^2 + x$ per x to check if the corresponding point lies on E . If so, we check if $x - \alpha$ is non-square via the Legendre symbol. On average, this requires four samplings of x and six Legendre symbols to find a suitable x_Q with $Q \in E(\mathbb{F}_{p^2})$, and, given that we can choose x_Q to be small, we can use fast scalar multiplication on x_Q .

In addition to computing α , smart sampling requires the Legendre symbol computation of $x^3 + Ax^2 + x$ per x . On average, we require two samplings of an x to find a suitable x_Q , hence saving four Legendre symbols in comparison to direct sampling. However, we can no longer choose x_Q small, which means that improved scalar multiplication for small x_Q is not available.

Finding Points with 2^f -Torsion *Not* Above $(0, 0)$. As shown in [26], we find a point P with full 2^f -torsion *not* above $(0, 0)$ by selecting a point on the curve with non-square x -coordinate. Non-squareness depends only on p , not on E , so a list of small non-square values can be precomputed. In this way, finding such a point P simply becomes finding the first value x_P in this list such that the point $(x_P, -)$ lies on $E(\mathbb{F}_{p^2})$, that is, $x_P^3 + Ax_P^2 + x_P$ is square. On average, this requires two samplings of x , hence two Legendre symbol computations.

5.2 General Improvements to Verification

In this section, we describe improvements to SQIsign verification and present new optimisations, decreasing the cost of the three main subroutines of verification.

Known Techniques from Literature. There are several state-of-the-art techniques in the literature on efficient implementations of elliptic curve or isogeny-based schemes that allow for general improvements to verification, but are not included in SQIsign (NIST). We implemented such methods, e.g., to improve scalar multiplication $P \mapsto [n]P$ and square roots. The details are described in the full version of our paper [10, Appendix A]. In particular, we use that $P \mapsto [n]P$ is faster when x_P is small.

Improving the Subroutine FindBasis. In SQIsign (NIST), to find a complete basis for $E[2^f]$ we are given a point $Q \in E[2^f]$ lying above $(0, 0)$ and need to find another point $P \in E(\mathbb{F}_{p^2})$ of order 2^f not lying above $(0, 0)$. We sample P directly using x_P non-square, as described above and demonstrated by [26], and in particular can choose x_P small. We then compute $P \leftarrow [\frac{p+1}{2^f}]P$ via fast scalar multiplication to complete the torsion basis (P, Q) .

Improved Strategies for Computelsogeny. Recall that Computelsogeny follows three steps in SQIsign (NIST): it first computes a 4-isogeny that may contain $(0, 0)$ in the kernel, and a 2-isogeny if f is odd, before entering an optimal strategy for computing the remaining chain of 4-isogenies. However, the first two steps include many costly doublings. We improve this by adding these first two steps in the optimal strategy. If f is even, this is straightforward, with a simple

check for $(0, 0)$ in the kernel in the first step. For odd f , we add the additional 2-isogeny in this first step.¹⁴ For simplicity of the implementation, we determine optimal strategies as in SIKE [22], thus we assume that only 4-isogenies are used.

Note that techniques for strategies with variable isogeny degrees are available from the literature on CSIDH implementations [9]. However, the performance difference is very small, hence our simplified approach appears to be preferable.

In addition to optimising 4-isogeny chains, we implemented optimised 3-isogeny chains from SIKE [22] for the computation of $\widehat{\varphi_{\text{chall}}}$ when $f < 128$.

5.3 To Push, or Not to Push¹⁵

In SQIsign (NIST), the point Q is pushed through φ so that we easily get the basis point above $(0, 0)$ on the image curve, and we can then use Theorem 2 to sample the second basis point P . Instead of pushing Q , one can also use Theorem 2 to efficiently sample this basis point Q above $(0, 0)$. Although pushing Q seems very efficient, for larger f we are pushing Q through increasingly larger isogeny chains, whereas sampling becomes increasingly more efficient as multiplication cost by $\frac{p+1}{2^f}$ decreases. Furthermore, sampling *both* P and Q allows us to use those points as an *implicit basis* for $E[2^f]$, even if their orders are multiples of 2^f , as described in more detail below. We observe experimentally that this makes sampling Q , instead of pushing Q , more efficient for $f > 128$.

Using Implicit Bases. Using Theorem 2, it is possible to find points P and Q efficiently so that both have full 2^f -torsion. The pair (P, Q) is not an *explicit basis* for $E[2^f]$, as the orders of these points are likely to be multiples of 2^f . However, instead of multiplying both points by the cofactor to find an explicit basis, we can use these points implicitly, as if they were a basis for $E[2^f]$. This allows us to compute $K = P + [s]Q$ first, and only then multiply K by the cofactor. This saves a full scalar multiplication by the cofactor $\frac{p+1}{2^f}$. We refer to such a pair (P, Q) as an *implicit basis* of $E[2^f]$. Algorithmically, implicit bases combine FindBasis and FindKernel into a single routine FindBasisAndKernel.

5.4 Improved Challenge for $f \geq \lambda$

Recall from Sect. 4.2 that when $f \geq \lambda$, we can simply set $D_{\text{chall}} = 2^\lambda$. This decreases the cost of FindBasis for the challenge computation considerably, as we can now use Theorem 2 to find a basis for $E[2^\lambda]$.

Improving FindBasis for the Challenge Isogeny When $f \geq \lambda$. We use Theorem 2 twice, first to find P not above $(0, 0)$ having full 2^f -torsion and then to find Q above $(0, 0)$ having full 2^f -torsion. We choose x_P and x_Q small such that faster scalar multiplication is available. We find the basis for $E[2^\lambda]$ by

¹⁴ In particular, we compute $R' = [2^{f-3}]K$ and $R = [2]R'$, a 4-isogeny with kernel $\langle R \rangle$, push R' through, and compute a 2-isogeny with kernel $\langle R' \rangle$.

¹⁵ –that is, the Q .

$P \leftarrow [\frac{p+1}{2^f}]P$ followed by $f - \lambda$ doublings, and $Q \leftarrow [\frac{p+1}{2^f}]Q$ followed by $f - \lambda$ doublings.¹⁶ Alternatively, if Q is pushed through isogenies, we can reuse $Q \leftarrow \varphi^{(n)}(Q^{(n)}) \in E[2^f]$ from the computation of the last step of φ_{resp} , so that we get a basis point for $E[2^\lambda]$ by $f - \lambda$ doublings of Q . Reusing this point Q also guarantees cyclicity of $\widehat{\varphi_{\text{chall}}} \circ \varphi_{\text{resp}}$.

Remark 7. For **SQIsign** without extension fields, obtaining $f \geq \lambda$ seems infeasible, hence the degree D of φ_{chall} is $2^f \cdot 3^g$. Nevertheless, some optimizations are possible in the computation of φ_{chall} in this case. **FindBasis** for $E[2^f \cdot 3^g]$ benefits from similar techniques as previously used in **SIDH/SIKE**, as we can apply known methods to improve generating a torsion basis for $E[3^g]$ coming from 3-descent [14, Sect. 3.3]. Such methods are an analogue to generating a basis for $E[2^f]$ as described in Theorem 2 and [26, Theorem 3].

6 Size-Speed Trade-Offs in SQIsign Signatures

The increase in f also enables several size-speed trade-offs by adding further information in the signature or by using uncompressed signatures. Some trade-offs were already present in earlier versions of **SQIsign** [18], however, by using large f and the improvements from Sect. 5, they become especially worthwhile.

We take a slightly different stance from previous work on **SQIsign** as for many use cases the main road block to using **SQIsign** is the efficiency of verification in cycles. In contrast, in several applications the precise size of a signature is less important as long as it is below a certain threshold.¹⁷ For example, many applications can handle the combined public key and signature size of RSA-2048 of 528 bytes, while **SQIsign** (NIST) features a combined size of only 241 bytes. In this section, we take the 528 bytes of RSA-2048 as a baseline, and explore size-speed trade-offs for **SQIsign** verification with data sizes up to this range.

We note that the larger signatures in this section encode the same information as standard **SQIsign** signatures, hence have no impact on the security.

6.1 Adding Seeds for the Torsion Basis in the Signature

We revisit an idea that was previously present in **SQIsign** verification [18] (but no longer in [8] or [19]), and highlight its particular merits whenever $f \geq \lambda$, as enabled by signing with extension fields. So far, we have assumed that completing or sampling a basis for $E[2^f]$ is done by deterministically sampling points. Recall from Sect. 5.1 that sampling x_P resp. x_Q (when not pushing Q) on average requires the computation of several Legendre symbols resp. square roots. We instead suggest using a seed to find x_P (when pushing Q) or x_P and x_Q (otherwise), which we include in the signature, so that the verifier saves all of the above cost for finding x_P , resp. x_Q . Finding these seeds adds negligible overhead

¹⁶ Algorithmically, this is faster than a single scalar multiplication by $2^{f-\lambda} \cdot \frac{p+1}{2^f}$.

¹⁷ See <https://blog.cloudflare.com/sizing-up-post-quantum-signatures/>.

for the signer, while verification performance improves. Signer and verifier are assumed to agree upon all precomputed values.

Seeding a Point *Not* Above $(0, 0)$. For x_P *not* above $(0, 0)$, we fix a large enough $k > 0$ and precompute the 2^k smallest values $u_j \in \mathbb{F}_p$ such that $u_j + i \in \mathbb{F}_{p^2}$ is non-square (where i is the same as in Sect. 5). During signing, we pick the smallest u_j such that $x_P = u_j + i$ is the x -coordinate of a point $P \in E(\mathbb{F}_{p^2})$, and add the index j to the signature as a seed for x_P . Theorem 2 ensures that any $P \in E(\mathbb{F}_{p^2})$ for non-square x_P is a point with full 2^f -torsion not above $(0, 0)$. This furthermore has the advantage of fast scalar multiplication for x_P as the x -coordinate is very small.

Seeding a Point *Above* $(0, 0)$. As noted above, when f is large, it is faster to deterministically compute a point of order 2^f above $(0, 0)$ than to push Q through φ . We propose a similar seed here for fixed large enough $k > 0$, using Theorem 2 and the “direct sampling” approach from Sect. 5.1. During signing, we pick the smallest $j \leq 2^k$ such that $x_Q = j$ is the x -coordinate of a point $Q \in E(\mathbb{F}_{p^2})$ and $x_Q - \alpha$ is non-square. We add $x_Q = j$ to the signature as seed.

Note that when using both seeding techniques, we do not explicitly compute $\lceil \frac{p+1}{2^f} \rceil P$ or $\lceil \frac{p+1}{2^f} \rceil Q$, but rather use the seeded points P and Q as an implicit basis, as described in Sect. 5.3.

Size of Seeds. Per seeded point, we add k bits to the signature size. Thus, we must balance k such that signatures are not becoming too large, while failure probabilities for not finding a suitable seed are small enough. In particular, seeding x_P resp. x_Q via direct sampling has a failure probability of $\frac{1}{2}$ resp. $\frac{3}{4}$ per precomputed value. For the sake of simplicity, we set $k = 8$ for both seeds, such that every seed can be encoded as a separate byte.¹⁸ This means that the failure rate for seeding Q is $(\frac{3}{4})^{256} \approx 2^{-106.25}$ for our choice, while for P it is 2^{-256} . Theoretically it is still possible that seeding failures occur. In such a case, we simply recompute KLPT. We furthermore include similar seeds for the torsion basis on E_A and E_2 , giving a size increase of $(n + 1) \cdot 2$ bytes.

The synergy with large f now becomes apparent. The larger f gets, the fewer blocks n are required, hence adding fewer seeds overall. For $f = 75$, the seeds require an additional 28 bytes when seeding both P and Q . For $f = 122, 140, 163, 195, 244$ this drops to 18, 16, 14, 12, and 10 additional bytes, respectively, to the overall signature size of 177 bytes for NIST Level I security.

Remark 8. Instead of using direct sampling for Q with failure probability $\frac{3}{4}$, we can reduce it to $\frac{1}{2}$ via “smart sampling” (see Sect. 5.1). However, this requires the verifier to compute α via a square root to set $x_Q = z\alpha$ with seeded z . We thus prefer direct sampling for seeded Q , which incurs no such extra cost.

¹⁸ Note that for equal failing rates the number of possible seeds for P can be chosen smaller than for Q , hence slightly decreasing the additional data sizes.

6.2 Uncompressed Signatures

In cases where f is very large, and hence the number of blocks is small, in certain cases it is worthwhile to replace the value s in the signature by the full x -coordinate of $K = P + [s]Q$. In essence, this is the uncompressed version of the SQIsign signature σ , and we thus refer to this variant as *uncompressed SQIsign*.

Speed of Uncompressed Signatures. Adding the precise kernel point K removes the need for both FindBasis and FindKernel, leaving ComputelSogeny as the sole remaining cost. This speed-up is significant, and leaves little room for improvement beyond optimizing the cost of computing isogenies. The cost of verification in this case is relatively constant, as computing an 2^e -isogeny given the kernels is only slightly affected by the size of f , as is visible in the black dashed line in Fig. 2. This makes uncompressed SQIsign an attractive alternative in cases where the signature size, up to a certain bound, is less relevant.

Size of Uncompressed Signatures. Per step, this increases the size from $\log(s) \approx f$ to $2 \cdot \log(p)$ bits, which is still relatively size efficient when f is close to $\log(p)$. For recomputing φ_{chall} , we take a slightly different approach than before. We add the Montgomery coefficient of E_1 to the signature, and seeds for a basis of $E[2^f]$. From this, the verifier can compute the kernel generator of φ_{chall} , and verify that the j -invariant of its codomain matches E_2 . Hence this adds $2 \cdot \log(p)$ bits for E_1 and two bytes for seeds to the signature, for a total of $(n + 1) \cdot (\log p / 4) + 2$ bytes.

For $f = 244$, this approach less than doubles the signature size from 177 bytes to 322 bytes for NIST Level I security, for $f = 145$, the signature becomes approximately 514 bytes, while for the current NIST Level I prime with $f = 75$, the size would become 898 bytes. When adding the public key size of 64 bytes, especially the first two cases still appear to be reasonable alternatives to RSA-2048's combined data size of 528 bytes.

Remark 9. Uncompressed signatures significantly simplify verification, as many functionalities required for compressed signatures are not necessary. Hence, this allows for a much more compact code base, which might be important for use cases featuring embedded devices with strict memory requirements.

7 Primes and Performance

In this section we show the performance of verification for varying f , using the optimisations from the previous sections. Further, we find specific primes with suitable f for $n = 4$ and $n = 7$, and report their signing performance using our SageMath implementation, comparing it with the current SQIsign (NIST) prime.




Fig. 2. Extended version of Fig. 1 showing the cost in \mathbb{F}_p -multiplications for verification at NIST Level I security, for varying f and $p^{(f)}$, averaged over 1024 runs per prime. In addition to SQIsign (NIST) in blue, it shows the performance of SQIsign (LWXZ) in red, our fastest compressed AprèsSQI variant in brown, and uncompressed AprèsSQI in black. (Color figure online)

7.1 Performance of Optimised Verification

To compare the verification performance of our optimised variants with compressed signatures to SQIsign (NIST) and SQIsign (LWXZ), we run benchmarks in the same setting as in Sect. 4.3. In particular, Fig. 2 shows the cost of verification for the NIST Level I primes $p^{(f)}$ for $50 \leq f \leq 250$. As before, we sample random public key curves and signatures σ of the correct form instead of using signatures generated by the SQIsign signing procedure.

For the sake of simplicity, Fig. 2 displays only the fastest compressed AprèsSQI variant, namely the version that does not push Q through isogenies and uses seeds to sample P and Q . This variant significantly outperforms both SQIsign (NIST) and SQIsign (LWXZ) already at $f = 75$, at the cost of slightly larger signatures. A detailed description and comparison of all four compressed variants is in the full version of our paper [10, Appendix C], which shows that our unseeded variants achieve similar large speed-ups with no increase in signature size. Lastly, the uncompressed variant achieves the fastest speed, although at a significant increase in signature size.

7.2 Finding Specific Primes

We now give two example primes, one prime optimal for 4-block verification, as well as the best we found for 7-block verification. The “quality” of a prime p is measured using the cost metric SIGNINGCOST_p defined in Sect. 3.3.

Optimal 4-Block Primes. For 4-block primes, taking $e = 975$ as a baseline, we need f bigger than 244. In other words, we are searching for primes of the form $p = 2^{244}N - 1$ where $N \in [2^4, 2^{12}]$ (accepting primes between 250 and 256 bits). This search space is quickly exhausted. For each prime of this form, we find the optimal torsion T to use, minimising $\text{SIGNINGCOST}_p(T)$. The prime with the lowest total cost in this metric, which we denote p_4 , is

$$p_4 = 2^{246} \cdot 3 \cdot 67 - 1$$

Balanced Primes. Additionally, we look for primes that get above the significant $f > 128$ line, while minimizing $\text{SIGNINGCOST}_p(T)$. To do this, we adopt the “sieve-and-boost” technique used to find the current `SQIsign` primes [8, Sect. 5.2.1]. However, instead of looking for divisors of $p^2 - 1$, we follow Theorem 1 and look for divisors of $\prod_{n=1}^k \Phi_n(p^2)/2$ to find a list of good candidate primes. This list is then sorted with respect to their signing cost according to SIGNINGCOST_p . The prime with the lowest signing cost we could find, which we call p_7 , is

$$p_7 = 2^{145} \cdot 3^9 \cdot 59^3 \cdot 311^3 \cdot 317^3 \cdot 503^3 - 1.$$

Remark 10. This method of searching for primes is optimised for looking for divisors of $p^2 - 1$, hence it might be suboptimal in the case of allowing torsion in higher extension fields. We leave it as future work to find methods which further take advantage of added flexibility in the prime search.

7.3 Performance for Specific Primes

We now compare the performance of the specific primes p_4 , p_7 , as well as the current NIST Level I prime p_{1973} used in `SQIsign` (NIST).

Signing Performance. We give a summary of the estimated signing costs in Table 1. For p_{1973} , we include the metric “Adjusted Cost”, which we compute as SIGNINGCOST with the isogeny computations scaling as $\sqrt{\ell} \log \ell$ to (rather optimistically) account for the benefit of $\sqrt{\ell}u$. Further, we ran our proof-of-concept SageMath implementation on the three primes, using SageMath 9.8, on a laptop with an Intel-Core i5-1038NG7 processor, averaged over five runs. An optimised C implementation will be orders of magnitude faster; we use these timings simply for comparison.

We note that the SIGNINGCOST -metric correctly predicts the ordering of the primes, though the performance difference is smaller than predicted. A possible explanation for this is that the SIGNINGCOST -metric ignores all overhead, such as quaternion operations, which roughly adds similar amounts of cost per prime.

Our implementation uses $\sqrt{\ell}u$ whenever the kernel generator is defined over \mathbb{F}_{p^2} and ℓ is bigger than a certain crossover point. This mainly benefits p_{1973} , as this prime only uses kernel generators defined over \mathbb{F}_{p^2} . The crossover point is experimentally found to be around $\ell > 300$ in our implementation, which

Table 1. Comparison between estimated cost of signing for three different primes.

p	largest $\ell \mid T$	largest $\mathbb{F}_{p^{2k}}$	$\text{SIGNINGCOST}_p(T)$	Adj. Cost	Timing
p_{1973}	1973	$k = 1$	8371.7	1956.5	11 m, 32 s
p_7	997	$k = 23$	4137.9	-	9 m, 20 s
p_4	2293	$k = 53$	9632.7	-	15 m, 52 s

Table 2. Comparison between verification cost for different variants and different primes, with cost given in terms of $10^3 \mathbb{F}_p$ -multiplications, using $\mathbf{S} = 0.8 \cdot \mathbf{M}$.

p	f	Implementation	Variant	Verif. cost	Sig. size
p_{1973}	75	SQlsign (NIST) [8]	-	500.4	177 B
		SQlsign (LWXZ) [26]	-	383.1	177 B
		AprèsSQI	unseeded	276.1	177 B
		AprèsSQI	seeded	226.8	195 B
p_7	145	AprèsSQI	unseeded	211.0	177 B
		AprèsSQI	seeded	178.6	193 B
		AprèsSQI	uncompressed	103.7	514 B
p_4	246	AprèsSQI	unseeded	185.2	177 B
		AprèsSQI	seeded	160.8	187 B
		AprèsSQI	uncompressed	112.2	322 B

is not optimal, compared to an optimised C implementation.¹⁹ Nevertheless, we believe that these timings, together with the cost metrics, provide sufficient evidence that extension field signing in an optimised implementation stays in the same order of magnitude for signing time as staying over \mathbb{F}_{p^2} .

Verification Performance. In Table 2, we summarise the performance of verification for p_{1973} , p_7 , and p_4 , both in terms of speed, and signature sizes.

Two highlights of this work lie in using p_7 , both with and without seeds, having (almost) the same signature sizes as the current SQlsign signatures, but achieving a speed-up of factor 2.37 resp. 2.80 in comparison to SQlsign (NIST) and 1.82 resp. 2.15 in comparison to SQlsign (LWXZ), using p_{1973} . Another interesting alternative is using uncompressed p_4 , at the cost of roughly double signature sizes, giving a speed-up of factor 4.46 in comparison to SQlsign (NIST) and 3.41 in comparison to SQlsign (LWXZ).

Remark 11. We analyse and optimise the cost of verification with respect to \mathbb{F}_p -operations. However, primes of the form $p = 2^f \cdot c - 1$ are considered to be particularly suitable for fast optimised finite field arithmetic, especially when

¹⁹ For instance, work by Adj, Chi-Domínguez, and Rodríguez-Henríquez [1] gives the crossover point at $\ell > 89$, although for isogenies defined over \mathbb{F}_p .

f is large [3]. Hence, we expect primes like p_4 to improve significantly more in comparison to p_{1973} in low-level field arithmetic, leading to a larger speed-up than predicted in Table 2. Furthermore, other low-level improvements, such as fast non-constant time GCD for inversions or Legendre symbols, will improve the performance of primes in terms of cycles, which is unaccounted for by our cost metric.

Acknowledgement. We thank Craig Costello for helpful suggestions and comments on an earlier version of this work. We thank the anonymous Eurocrypt 2024 reviewers for their constructive feedback.

References

1. Adj, G., Chi-Domínguez, J.-J., Rodríguez-Henríquez, F.: Karatsuba-based square-root Vélu’s formulas applied to two isogeny-based protocols. *J. Cryptogr. Eng.* **13**(1), 89–106 (2023). <https://doi.org/10.1007/s13389-022-00293-y>
2. Auer, R., Top, J.: Legendre elliptic curves over finite fields. *J. Number Theor.* **95**(2), 303–312 (2002). ISSN 0022-314X. <https://doi.org/10.1006/jnth.2001.2760>. <https://www.sciencedirect.com/science/article/pii/S0022314X0192760X>
3. Bajard, J.-C., Duquesne, S.: Montgomery-friendly primes and applications to cryptography. *J. Cryptogr. Eng.* **11**(4), 399–415 (2021). <https://doi.org/10.1007/s13389-021-00260-z>
4. Banegas, G., Gilchrist, V., Le Dévéhat, A., Smith, B.: Fast and Frobenius: rational isogeny evaluation over finite fields. In: Aly, A., Tibouchi, M. (eds.) *Progress in Cryptology, LATINCRYPT 2023*. LNCS, vol. 14168, pp. 129–148. Springer, Cham (2023). https://doi.org/10.1007/978-3-031-44469-2_7
5. Bernstein, D.J., De Feo, L., Leroux, A., Smith, B.: Faster computation of isogenies of large prime degree. In: *Open Book Series*, vol. 4, no. 1, pp. 39–55 (2020)
6. Biasse, J.-F., Jao, D., Sankar, A.: A quantum algorithm for computing isogenies between supersingular elliptic curves. In: Meier, W., Mukhopadhyay, D. (eds.) *INDOCRYPT 2014*. LNCS, vol. 8885, pp. 428–442. Springer, Cham (2014). https://doi.org/10.1007/978-3-319-13039-2_25
7. Bruno, G., et al.: Cryptographic Smooth Neighbors. *IACR Cryptol. ePrint Arch.*, p. 1439 (2022). <https://eprint.iacr.org/2022/1439>
8. Chavez-Saab, J., et al.: SQIsign: algorithm specifications and supporting documentation (2023). National Institute of Standards and Technology. <https://csrc.nist.gov/csrc/media/Projects/pqc-dig-sig/documents/round-1/spec-files/sqisign-spec-web.pdf>
9. Chi-Domínguez, J.-J., Rodríguez-Henríquez, F.: Optimal strategies for CSIDH. *Adv. Math. Commun.* **16**(2), 383–411 (2022)
10. Santos, M.C.-R., Eriksen, J.K., Meyer, M., Reijnders, K.: AprèsSQI: extra fast verification for SQIsign using extension-field signing. *Cryptology ePrint Archive, Paper 2023/1559* (2023). <https://eprint.iacr.org/2023/1559>
11. Costello, C.: B-SIDH: supersingular isogeny Diffie-Hellman using twisted torsion. In: Moriai, S., Wang, H. (eds.) *ASIACRYPT 2020, Part II*. LNCS, vol. 12492, pp. 440–463. Springer, Cham (2020). https://doi.org/10.1007/978-3-030-64834-3_15
12. Costello, C.: Computing supersingular isogenies on Kummer surfaces. In: Peyrin, T., Galbraith, S. (eds.) *ASIACRYPT 2018, Part III*. LNCS, vol. 11274, pp. 428–456. Springer, Cham (2018). https://doi.org/10.1007/978-3-030-03332-3_16

13. Costello, C., Hisil, H.: A simple and compact algorithm for SIDH with arbitrary degree isogenies. In: Takagi, T., Peyrin, T. (eds.) ASIACRYPT 2017, Part II. LNCS, vol. 10625, pp. 303–329. Springer, Cham (2017). https://doi.org/10.1007/978-3-319-70697-9_11
14. Costello, C., Jao, D., Longa, P., Naehrig, M., Renes, J., Urbanik, D.: Efficient compression of SIDH public keys. In: Coron, J.-S., Nielsen, J.B. (eds.) EUROCRYPT 2017, Part I. LNCS, vol. 10210, pp. 679–706. Springer, Cham (2017). https://doi.org/10.1007/978-3-319-56620-7_24
15. Costello, C., Meyer, M., Naehrig, M.: Sieving for twin smooth integers with solutions to the Prouhet-Tarry-Escott problem. In: Canteaut, A., Standaert, F.-X. (eds.) EUROCRYPT 2021, Part I. LNCS, vol. 12696, pp. 272–301. Springer, Cham (2021). https://doi.org/10.1007/978-3-030-77870-5_10
16. Dartois, P., Leroux, A., Robert, D., Wesolowski, B.: SQISignHD: new dimensions in cryptography. IACR Cryptol. ePrint Arch., p. 436 (2023). <https://eprint.iacr.org/2023/436>
17. De Feo, L., Jao, D., Plût, J.: Towards quantum-resistant cryptosystems from supersingular elliptic curve isogenies. J. Math. Cryptol. **8**(3), 209–247 (2014). <https://doi.org/10.1515/jmc-2012-0015>
18. De Feo, L., Kohel, D., Leroux, A., Petit, C., Wesolowski, B.: SQISign: compact post-quantum signatures from quaternions and isogenies. In: Moriai, S., Wang, H. (eds.) ASIACRYPT 2020, Part I. LNCS, vol. 12491, pp. 64–93. Springer, Cham (2020). https://doi.org/10.1007/978-3-030-64837-4_3
19. De Feo, L., Leroux, A., Longa, P., Wesolowski, B.: New algorithms for the deuring correspondence. In: Hazay, C., Stam, M. (eds.) Advances in Cryptology, EUROCRYPT 2023. LNCS, vol. 14008, pp. 659–690. Springer, Cham (2023). https://doi.org/10.1007/978-3-031-30589-4_23
20. Delfs, C., Galbraith, S.D.: Computing isogenies between supersingular elliptic curves over \mathbb{F}_p . Des. Codes Crypt. **78**, 425–440 (2016)
21. Eriksen, J.K., Panny, L., Sotáková, J., Veroni, M.: Deuring for the people: supersingular elliptic curves with prescribed endomorphism ring in general characteristic. IACR Cryptol. ePrint Arch., p. 106 (2023). <https://eprint.iacr.org/2023/106>
22. Jao, D., R., et al.: SIKE. Technical report, National Institute of Standards and Technology (2022). <https://csrc.nist.gov/Projects/post-quantum-cryptography/round-4-submissions>
23. Johnson, D., Menezes, A., Vanstone, S.A.: The elliptic curve digital signature algorithm (ECDSA). Int. J. Inf. Sec. **1**(1), 36–63 (2001). <https://doi.org/10.1007/s102070100002>
24. Josefsson, S., Liusvaara, I.: Edwards-curve digital signature algorithm (EdDSA). RFC: 8032, pp. 1–60 (2017). <https://doi.org/10.17487/RFC8032>
25. Kohel, D., Lauter, K., Petit, C., Tignol, J.-P.: On the quaternion-isogeny path problem. LMS J. Comput. Math. **17**(A), 418–432 (2014)
26. Lin, K., Wang, W., Xu, Z., Zhao, C.-A.: A faster software implementation of SQISign. Cryptology ePrint Archive, Paper 2023/753 (2023). <https://eprint.iacr.org/2023/753>
27. Meyer, M., Reith, S.: A faster way to the CSIDH. In: Chakraborty, D., Iwata, T. (eds.) INDOCRYPT 2018. LNCS, vol. 11356, pp. 137–152. Springer, Cham (2018). https://doi.org/10.1007/978-3-030-05378-9_8
28. Page, A., Wesolowski, B.: The supersingular endomorphism ring and one endomorphism problems are equivalent. CoRR abs/2309.10432. arXiv [arXiv:2309.10432](https://arxiv.org/abs/2309.10432) (2023). <https://doi.org/10.48550/arXiv.2309.10432>

29. Shor, P.W.: Polynomial-time algorithms for prime factorization and discrete logarithms on a quantum computer. *SIAM Rev.* **41**(2), 303–332 (1999)
30. Shoup, V.: Efficient computation of minimal polynomials in algebraic extensions of finite fields. In: *Proceedings of the 1999 International Symposium on Symbolic and Algebraic Computation*, pp. 53–58 (1999)
31. Silverman, J.H.: *The Arithmetic of Elliptic Curves*, vol. 106. Springer, New York (2009). <https://doi.org/10.1007/978-1-4757-1920-8>
32. National Institute of Standards and Technology (NIST): Call for Additional Digital Signature Schemes for the Post-Quantum Cryptography Standardization Process (2022). <https://csrc.nist.gov/csrc/media/Projects/pqc-dig-sig/documents/call-for-proposals-dig-sig-sept-2022.pdf>
33. Tsukazaki, K.: Explicit isogenies of elliptic curves. Ph.D. thesis, University of Warwick (2013)
34. Vélú, J.: Isogénies entre courbes elliptiques. *Comptes-Rendus de l’Académie des Sciences* **273**, 238–241 (1971)
35. Voight, J.: *Quaternion Algebras*. GTM, vol. 288. Springer, Cham (2021). <https://doi.org/10.1007/978-3-030-56694-4>
36. Wesolowski, B.: The supersingular isogeny path and endomorphism ring problems are equivalent. In: *2021 IEEE 62nd Annual Symposium on Foundations of Computer Science (FOCS)*, pp. 1100–1111. IEEE (2022)